

Hermes CIO Build — Time Estimate

Polished product, including tests, QA and deploy.

Spec	<code>docs/build-28th-may.md</code>
Goal	Add the Hermes CIO layer to DJG Advisory and replace the OpenClaw agent (Mac mini) with Hermes as the single CIO agent.
Dated	29 May 2026

SCOPE

- Steno Mirror is kept — tab and pipeline stay. The earlier removal and alt-coin audit task is dropped.
- The Hermes agent layer runs on GPT 5.5 (OpenAI), not Anthropic. A new OpenAI client and tool-use loop is required; the existing Anthropic helper is a reference pattern only.

ALREADY IN PLACE

- 7-factor scenario engine — `risk_on_off`, `growth`, `inflation`, `short_rates`, `liquidity`, `dollar`, `oil` (`correlations.py`). These are exactly the Risk Budget inputs.
- Live macro to $-1..+1$ auto-fill via rolling z-scores (`scenario_inputs.py`). Roughly 80% of the Risk Budget math.
- Backtest service with price history and forward-return machinery (`backtest.py`). Foundation for What-If.
- Telegram client, scheduler, daily card and regime alerts (`monitor.py`). The pattern for the weekly message.
- `launchd` deploy plist (`deploy/com.june.dashboard.plist`). Mac-mini deploy infrastructure.

ESTIMATE BY COMPONENT

#	COMPONENT	NORMAL	AI-ENABLED	NOTES
A	Nav cleanup (Allocation / Liquidity / Cycle, keep Steno)	1–2	0.5–1	Minor only
B	Risk Budget engine (0–100) and 5-band stance map	8–12	3–4	Inputs exist; weighting and tests
C	Hermes CIO View card (state aggregation, frontend)	12–16	4–6	Summary plus top-of-page card
D	What-If Outcome engine, chart and form	14–20	5–7	CAGR, max drawdown, monthly extremes
E	Framework Portfolio backtest	24–40	10–16	Highest risk — see below
F	Hermes agent on GPT 5.5 (OpenAI client, 9 tools, Q&A)	19–28	7–11	Net-new OpenAI integration
G	Weekly Telegram message with what-changed diff	8–12	3–4	Needs prior-week state diff
H	MIT overlay and SLR / eSLR note	3–5	1–2	Small qualitative fields
I	Integration, QA, deploy to Mac mini	16–24	6–10	Replace OpenClaw
Total		105–159h	40–61h	

BOTTOM LINE

NORMAL CODING

~129h

105–159h · about 3–4 weeks

AI-ENABLED CODING

~50h

40–61h · about 1.5–2 weeks

KEY RISK

Item E, Framework Portfolio, is the spec's "important one" and the biggest unknown. The allocation engine uses the live correlation matrix, not point-in-time data. A faithful backtest needs either rolling point-in-time matrices (expensive) or a simplified regime-to-allocation map. Decide the fidelity up front — it alone swings the total by about 15h normal / 6h AI.

